MoCHII User Guide

Abstract

How to build and run MoCHII (MOnte Carlo for H II regions): requirements, the complete input-parameter reference, output-file formats, worked examples keyed to the regression tests, and the runbook for adding an ion. The physics and validation are documented separately in `docs/MoCHII_physics.pdf`; the atomic-data fitting pipeline in `docs/MoCHII_fitting.pdf`.

Contents

1. Requirements and build

- Intel oneAPI MPI Fortran (`mpiifort`/`mpiifx`; GNU `mpif90` supported by the Makefile), HDF5 (default prefix `/data/opt/hdf5_intel`; override with `make HDF5_PREFIX=...`), CFITSIO.
- The SEDust dust-emission library, self-contained under `SEDust/` in this tree. Build it once with `cd SEDust/sed && ./build_lib.sh` (→ `SEDust/sed/lib/libsedust.a` and its `.mod` files, which the Makefile links). The SEDust *data* directory is read at run time and resolved automatically relative to the executable (Section 2.10).
- Python 3 with numpy/scipy (h5py, astropy, matplotlib, PyNeb for the test and pipeline scripts).

Build:

```
cd MoCHII
(cd SEDust/sed && ./build_lib.sh)   # on a fresh checkout
make            # -> MoCHII.x   (always a full clean rebuild)
make DEBUG=1    # bounds/traceback build
```

The SEDust library only needs the one-time `build_lib.sh` step; afterward `make` links the in-tree SEDust library. The Makefile intentionally performs a full rebuild every time (the MoCafe policy): there is no inter-module dependency tracking, and stale `.mod` files after editing `define.f90` otherwise cause runtime namelist errors.
Run:

```
mpirun -np 8 ./MoCHII.x input.in
```

MoCHII is MPI-only; the octree and all leaf arrays live in MPI-3 shared memory (one copy on each node).

2. Input parameters

All input is one `&parameters` namelist assigning fields of `par`. Parameters not listed keep the defaults shown. The switches compose: `gas_niter = 0` gives a single rates-only pass; `gas_niter > 0` iterates the ionization; `solve_te` adds the thermal balance; `use_metals` the trace-metal registry; `diffuse_field` the explicit diffuse field (forces case A); `refine_front` the I-front re-refinement; `ion_add_dust` the dust coupling.

2.1 Photons and RNG

parameter	default	meaning
`no_photons`	1e5	source packets in each iteration
`no_print`	1e7	progress-print stride
`iseed`	0	RNG seed (0 = from clock/pid)
`launch_sequence`	'random'	quasi-random launch: `'sobol'` = Owen-scrambled Sobol coordinates for the stellar (and diffuse) packet launch, indexed by the global photon number (MPI-count-independent launch set); transport/scattering stay pseudo-random; prefer `nphotons` = 2^m. `'random'` is the Mersenne Twister default
`qmc_seed`	12345	scramble key (`launch_sequence='sobol'`); fixed across the gas iterations; change it for independent replicates (their scatter is the RQMC error estimate)

2.2 Grid and dust density

parameter	default	meaning
`grid_type`	'car'	`'amr'` (octree) or `'car'` (single-level Cartesian, leaf list in raster order; no tree and no geometry arrays in memory — cell centers and half-widths are computed analytically from the raster index; `refine_front` unavailable). `'uniform'` is accepted as an alias for `'car'`
`car_walk`	'dda'	car-grid traversal: `'dda'` (default) = dedicated incremental Amanatides–Woo walks ($t_{\max}/t_{\Delta}$ per axis, comparisons and additions only — the fastest car walk); `'shared'` = the octree walks with integer index arithmetic replacing the neighbor table, a verification mode that is bit-identical to `'amr'`
`amr_file`	''	generic AMR leaf file (FITS/HDF5/text): columns x, y, z, level, nH [, metallicity, xHI, ndust]. If empty, a `'car'` grid is built from the namelist instead (`'amr'` requires the file)
`nx, ny, nz`	1, 1, 11	`'car'`-from-namelist cell counts (cubic cells required: $2x_{\max}/n_x = 2y_{\max}/n_y = 2z_{\max}/n_z$)
`xmax, ymax, zmax`	1, 1, 1	`'car'`-from-namelist box half-extents [`distance_unit`]; box is $[-x_{\max}, x_{\max}]$ etc.
`nH_const`	−1	if ≥ 0, override every leaf's $n_{\rm H}$ with this uniform value [cm^{-3}] (both `'car'` and `'amr'`); < 0 keeps the file/column density
`density_file`	''	read the leaf $n_{\rm H}$ from a uniform 3D density cube onto a `'car'` grid (FITS 3D image, or the MoCHII HDF5 image format; NAXIS$1/2/3 = n_x/n_y/n_z$, values $n_{\rm H}$ [cm^{-3}]). $n_x/n_y/n_z$ come from the file, the box is `xmax/ymax/zmax`; an alternative to `amr_file`. `reduce_factor`> 1 downsamples; a `rmax/rmin` cut can mask the cube to a sphere
`rmax, rmin`	−999, 0	spherical density cut on leaf centers (both grids): $n_{\rm H} = 0$ for $r > r_{\max}$ (`rmax`> 0) and for $r < r_{\min}$ (`rmin`> 0; a shell); radii in `distance_unit`
`distance_unit`	'kpc'	'kpc'/'pc'/'au'/'' (code units)

parameter	default	meaning
`dust_model`	'global_dgr'	leaf grey opacity: `'none'`, `'global_dgr'` ($n_{\rm H} c_{\rm ext}$·DGR), `'from_file'` (ndust column), `'laursen09'` (static xHI column), `'laursen09_live'` (computed $x_{\rm HI}$, refreshed each iteration)
`cext_dust`	1.6059e-21	reference extinction per H [cm^2] at `lambda_ref`
`DGR`, `Z_global`, `Z_ref`	1e-2, 0.0134, 0.0134	dust-model scalings
`f_ion_dust`	0.01	laursen09 survival of dust in ionized gas
`f_pah`, `f_ion_pah`	0, 0	PAH share of the reference extinction and its own ionized-gas survival (laursen09_live)

2.3 Ionizing/FUV band and the source

parameter	default	meaning
use_ion_band	F	must be T (the MoCHII transport)
nnu_ion	16	ionizing-band bins; gates use 32. With ion_align_edges on (the default) the bins are log-uniform within each threshold-bounded segment; off, they are a single log grid
eion_min, eion_max	13.598, 100	ionizing band edges [eV]
ion_align_edges	T	pin the ionizing-band bin edges onto the ionization thresholds (H I 13.6, He I 24.6, He II 54.4 eV plus the active metal edges) so no bin straddles one; segments between thresholds are log-uniform, and if the thresholds outnumber the bins the lowest-priority metal edges are dropped (H/He always kept), so it is safe at any nnu_ion. Removes the $\sim$18% He-ionizing-photon excess of a coarse log grid crossing the 24.6 eV edge. Set F for the legacy single log grid
add_fuv	F	FUV extension: extra log bins on [efuv_min, eion_min] below the unchanged ionizing bins; transports FUV grain heating and FUV photoionization of low-threshold metals (Mg I, C I, S I, Fe I); luminosity stays the IONIZING luminosity ($Q_{\rm H}$ preserved, the FUV part rides on top)
efuv_min	6.0	FUV lower edge [eV]
nnu_fuv	8	FUV bins (16 recommended with metals)
tstar	–999	Planck source temperature [K]
ion_spectrum	”	global spectrum file (column 1 plus one value column); its column units are set by spectrum_type (§2.5)
spectrum_type	’shape’	column units of every spectrum-file slot (ion_spectrum, src_spectrum_file, ext_spectrum): 'shape' (arbitrary, normalized to the scale — legacy) or a physical type 'per_ev'/'per_hz'/'per_ang'/'per_um' (absolute); see §2.5
luminosity	unset	ionizing BAND luminosity [erg/s] of the (single) source; for a 'shape'/Planck source unset maps to 1.0 (legacy), for a physical-type file unset means DERIVE it from the file integral; with nsource$\geq$ 2 it becomes the sum of src_lum (or the equal-split total when those are unset)
xs_point, ys_point, zs_point	0	single-source position (code units, recentered box; nsource=1)

Multiple internal point sources (the scalars above are the nsource=1 case):

parameter	default	meaning
`nsource`	1	number of internal point sources (0–16); 1 keeps the legacy scalars above, ≥ 2 uses the `src_*` arrays
`src_x`, `src_y`, `src_z`	0	positions of sources 1..`nsource` (code units)
`src_lum`	unset	ionizing luminosities [erg/s] of the sources; set all or none (none = equal split of `luminosity`); a partial set aborts
`src_tstar`	−999	Planck temperature [K] for each source (multi-source runs only; a single source uses `tstar`)
`src_spectrum_file`	”	one multi-column spectrum file for all sources: column 1 then one value column for each source (units set by `spectrum_type`, §2.5); each source is independently normalized so its ionizing segment carries `src_lum`(*i*) (a physical type with `src_lum` unset is derived from that column); overrides `src_tstar`; too few columns abort, extra columns are ignored with a note

External isotropic field (composes with the internal sources; ON when `ext_intensity`> 0, an ISRF preset is named, or a physical-type `ext_spectrum` file is given):

parameter	default	meaning
`ext_intensity`	−1	band-integrated mean intensity J [erg s^{-1} cm^{-2} sr^{-1}] of the external field; the field is ON when > 0 and adds to the internal sources; entering ionizing power $L = \pi J A_{\rm surf}$. With a physical-type `ext_spectrum` file or an ISRF preset it is optional: unset = the absolute J is DERIVED from the file/preset, set = the field is RESCALED to this band J (a note is printed)
`ext_geometry`	'rec'	illuminating surface: `'rec'` (box faces, face-area-weighted — correct for a non-cubic box) or `'sph'` (bounding sphere of radius `rmax` about the box center; needs `rmax`> 0)
`ext_tstar`	–999	external-field Planck temperature [K]
`ext_spectrum`	"	external-field spectrum: a 2-column file (units set by `spectrum_type`, §2.5; column 2 is an intensity density J) OR a reserved ISRF preset name `'draine'`/`'habing'`/`'mathis'` (FUV-only, absolute, needs `add_fuv`); priority `ext_spectrum` > `ext_tstar` > the global spectrum
`ext_scale`	1.0	dimensionless multiplier on an ISRF preset (ignored when `ext_intensity`> 0, which rescales the preset to a target band J)
`source_geometry`	'point'	legacy external-only shorthand: `'external'`, `'external_rec'`, or `'external_sph'` forces `nsource=0` with the external field on (needs `ext_intensity`> 0); `'point'` keeps the composable form; `'slab'` is the plane-parallel slab illumination below

2.4 Plane-parallel slab illumination

With `source_geometry = 'slab'` (which requires `xy_periodic`, §2.2) the box is an *xy-periodic slab*: the medium varies in z (uniform in x, y for a true 1D column with `nx = ny = 1`, or an arbitrary 3D tile with `nx, ny` > 1 repeated periodically), and light enters through the top ($+z$) and/or bottom ($-z$) face. Each face is either a collimated *beam* at incidence angle θ or an *isotropic* (Lambert) field, with its own strength. The single physical normalization is the flux F_z crossing the (horizontal) illuminated face: for a beam $F_z = \mu F_n$ ($\mu = \cos\theta$, F_n the flux normal to the beam), for the isotropic mode $F_z = \pi I$ (I the specific intensity). F_z sets the packet weight through $L = F_z A_{\rm zface}$; every per-volume result is independent of the tile area. Point-observer peel-off is disabled under `xy_periodic` (the periodic images have no finite-distance observer); the emergent observable is the boundary intensity $I(\mu)$ (§3.3).

parameter	default	meaning
`slab_faces`	'top'	which z-face(s) are lit: `'top'`, `'bottom'`, or `'both'`
`slab_top_mode`, `slab_bot_mode`	'beam'	per face: `'beam'` (collimated) or `'isotropic'` (Lambert)
`slab_top_theta`, `slab_bot_theta`	0	beam incidence angle from the inward normal [deg]; 0 is normal incidence
`slab_top_phi`, `slab_bot_phi`	0	beam azimuth [deg]
`slab_top_source`, `slab_bot_source`	0	per-face source strength (band-integrated): the specific intensity I [erg s^{-1} cm^{-2} sr^{-1}] for `'isotropic'`, or the beam flux F_n [erg s^{-1} cm^{-2}] for `'beam'`
`slab_nmu`	20	number of μ-bins of the emergent $I(\mu)$ output

The spectrum is the global one (`tstar` / `ion_spectrum`). Both z-faces always escape (no reflecting boundary). `launch_sequence = 'sobol'` is supported (both modes). With dust scattering (`ion_dust_scatter`) the scattered/reflected component is included in the boundary $I(\mu)$.

2.5 Spectrum-file formats

A single parameter, `spectrum_type` (default `'shape'`), fixes the column units of *every* spectrum-file slot (`ion_spectrum`, `src_spectrum_file`, `ext_spectrum`). The type sets the units; the slot sets the meaning: an internal-source file gives a luminosity density L_X, the external-field file gives an intensity density J_X. Column 1 and the value columns are:

`spectrum_type`	column 1	internal (cols 2+)	external (col 2)
`'shape'`	E [eV] ↑	arbitrary per eV (shape of L_E)	arbitrary per eV
`'per_ev'`	E [eV]	L_E [erg s^{-1} eV^{-1}]	J_E
`'per_hz'`	ν [Hz]	L_ν [erg s^{-1} Hz^{-1}]	J_ν
`'per_ang'`	λ [Å]	L_λ [erg s^{-1} Å^{-1}]	J_λ
`'per_um'`	λ [μm]	L_λ [erg s^{-1} μm^{-1}]	J_λ

The external intensity density J_X carries the same spectral unit as the internal L_X times an extra cm^{-2} sr^{-1} (a mean intensity). All files are plain text with `'#'` comment lines; values outside the tabulated range are zero. `src_spectrum_file` is one multi-column file: column 1 shared by all sources, then one value column for each source (extra columns beyond `nsource` are ignored with a note, too few abort).

The abscissa is a sample grid, not bin centers. Column 1 lists the points at which the spectral density is *sampled*: the file describes a continuous function, defined between the points by linear interpolation, and no bin widths are attached to the rows. MoCHII integrates each of its own band bins on a 32-point sub-grid of the bin, interpolating the table at every sub-point (zero outside the tabulated range), so the file grid and the band grid are independent. The spacing is arbitrary — nonuniform grids are fine; place points densely where the spectrum varies rapidly, since only linear variation is captured between adjacent points. Abscissa values must be distinct (no duplicates); a `'shape'` file must be ascending in E, while the physical types accept any row order (the table is sorted internally after the unit conversion, so a wavelength table may be given wavelength-ascending).

Absolute vs. shape normalization. A physical type (`'per_ev'`/`'per_hz'`/`'per_ang'`/`'per_um'`) is *absolute*: the bin luminosities follow the tabulated integral directly. If the matching scale (`luminosity` / `src_lum`(i) / `ext_intensity`) is unset it is DERIVED from the file — each source from its column's ionizing-segment integral, the external field from its band-integrated J — and reported in the log. If a scale is set, the spectrum is RESCALED to it (a note is printed) and the file then acts only as a shape. With `add_fuv` an absolute spectrum fills the FUV bins from the same table, with no separate normalization. `'shape'` keeps the legacy behavior: the shape is normalized to the scale in every case. A spectrum with no luminosity inside the band aborts, as does rescaling a file whose ionizing-segment integral is zero.

`'shape'` is per eV, not variable-agnostic. `'shape'` reads its value columns as the shape of the per-eV density L_E on the E [eV] grid of column 1. A table of F_λ (or F_ν) values re-gridded onto E under `'shape'` therefore omits the Jacobian and is wrong by a factor E^{-2} (the $\mathrm{d}\lambda/\mathrm{d}E$ conversion). To use a wavelength- or frequency-tabulated shape, give it with its matching physical type and set the scale explicitly — the type conversion then applies the Jacobian and the scale RESCALES the file to your target, so the file acts as a shape with the correct units. For example, a wavelength shape normalized to 10^{49} ionizing photons' worth of luminosity:

```
par%spectrum_type = 'per_ang'
par%ion_spectrum  = 'wave_shape.txt'
par%luminosity    = 3.0e38
```

A small `'per_ang'` source file for two point sources (λ ascending; the second source is FUV-faint):

```
# lambda[A]   L_lam,1        L_lam,2
   300.0      7.80e34        3.4e33
   500.0      6.05e34        1.2e33
   912.0      4.10e34        0.0
```

2.6 Initial state and iteration

parameter	default	meaning
`xHI_init`	1.0	initial $n_{\rm HI}/n_{\rm H}$ (an `xHI` file column wins); an ionized start (1e-3) converges much faster than neutral
`xHeI_init`, `xHeII_init`	1, 0	initial He fractions
`He_abund`	0.1	$n_{\rm He}/n_{\rm H}$
`gas_niter`	0	iterations (0 = a single rates-only pass)
`gas_tol`	1e-3	convergence tolerance on $x_{\rm HII}$ (measure set by `conv_crit`)
`conv_crit`	'cell'	convergence measure: 'cell' (cell maxima; stalls at the front-cell noise floor) or 'vol' (volume-integrated L_1; noise-robust, recommended). See §2.7
`ion_relax`	1.0	under-relaxation on x (< 1 damps front oscillation)
`case_ab`	'B'	case A/B recombination in the ionization balance, i.e. how the ionizing *continuum* is treated ('A' is forced by `diffuse_field`, which transports the ground-state recombination photons)
`line_case`	'B'	case of the SH95 H I/He II recombination-*line* tables. This is set by the optical depth of the Lyman *lines*, a different question from `case_ab`: a nebula whose diffuse continuum is transported still traps its Lyman lines and radiates case-B Balmer lines, so 'B' is the physical default for H II regions. 'A' selects the case-A tables (Lyman-line-thin gas); 'follow' ties the line case to `case_ab`
`require_convergence`	F	if the gas iteration has not converged at `gas_niter`, print a warning; when T, abort before writing output. The convergence status is recorded in the `_rates` header either way (`CONVERGD`, `NITERDON`, `FINALDX`, `FINALDTE`)

2.7 Convergence of the gas iteration

The gas iteration (`gas_niter` sweeps of transport → rate integrals → leaf solve → opacity refill) stops when the state stops changing between sweeps. Each sweep compares a *change measure* against a tolerance:

- ionization only (`te_fixed`, no thermal solve): converged when the $x_{\rm HII}$ change $<$ `gas_tol`;
- coupled (`solve_te`): converged when the $x_{\rm HII}$ change $<$ `gas_tol` *and* the T_e change $<$ `gas_tol_te`.

Two measures (`conv_crit`). Both are computed every sweep; `conv_crit` selects which one gates the loop.

`'cell'` the cell maxima $\max_{\rm leaf}|\Delta x_{\rm HII}|$ and $\max_{\rm leaf}|\Delta T_e|/T_e$. Simple, but the ionization-front cells jitter with Monte Carlo noise and never settle, so $\max|\Delta x_{\rm HII}|$ *stalls* at the front-cell noise floor ($\sim$2–4$\times10^{-2}$ at 4×10^7 packets) and a smaller `gas_tol` is unreachable. This is the historical default (kept so the recorded gates reproduce).

'vol' the volume-integrated L_1 changes

$$\frac{\sum_{\rm leaf} V\,|\Delta x_{\rm HII}|}{\sum_{\rm leaf} V\,x_{\rm HII}} < \texttt{gas\_tol}, \qquad \frac{\sum_{\rm leaf} V\,n_e\,|\Delta T_e|}{\sum_{\rm leaf} V\,n_e\,T_e} < \texttt{gas\_tol\_te}.$$

Front jitter occupies little volume (a few thousand front cells against $\sim 10^5$ ionized ones), so it enters the norm suppressed by that ratio; the n_e weight drops the vacuum and `te_min`-pinned no-heating cells that otherwise dominate $\max|\Delta T_e|$ in FUV/PDR runs. These measures decay smoothly — the recommended production setting.

The Monte Carlo floor. Even the volume measures do not reach zero: they floor at a Monte Carlo level set by the packet count, scaling as $1/\sqrt{N_{\rm packets}}$ ($\sim 3\times10^{-3}$ in x and $\sim 7\times10^{-3}$ in T_e at 8×10^6 packets). Reaching the floor *is* convergence — iterating further leaves the solution unchanged. So **set `gas_tol`/`gas_tol_te` just above the floor for the packet count in use**: a tolerance below the floor can never be met and the loop runs to `gas_niter` and warns. Rules of thumb (`conv_crit='vol'`): $\sim 5\times10^{-3}/10^{-2}$ at $\sim 10^7$ packets, loosened to $\sim 10^{-2}/2\times10^{-2}$ at $\sim 10^6$ (the reduced-count `examples/` use the looser values for this reason).

When it does not converge. Hitting `gas_niter` without meeting the test is not fatal: the state is still written and a warning is printed on rank 0. Set `require_convergence = T` to abort before the output instead. Either way the `_rates` header records the outcome in `CONVERGD` (0/1), `NITERDON` (sweeps run), `FINALDX`, and `FINALDTE`. The derivation and the gate measurements are in `docs/MoCHII_physics.pdf`.

2.8 Thermal balance and metals

parameter	default	meaning
`solve_te`	F	bisection on T_e coupled with the ionization
`te_fixed`	1e4	fixed T_e [K] when `solve_te` is off
`te_min`, `te_max`	1e3, 5e4	bisection bracket [K]
`gas_tol_te`	1e-3	convergence on $\max\|\Delta T_e\|/T_e$
`atomic_dir`	'data/atomic'	coefficient files (relative to the run directory)
`use_metals`	F	registry cascade + metal line cooling
`ci_model`	'voronov'	collisional ionization: 'voronov' (default) or 'dere_hybrid' (Voronov × the Dere 2007/Voronov ratio; matters only in shielded low-ionization zones)
`abund_C/N/O/Ne/S`	HII20 values	$n(\mathrm{X})/n(\mathrm{H})$; an element is active when > 0
`abund_Ar/Mg/Fe`	0	further registry elements (gas-phase values; Mg/Fe are heavily depleted)
`abund_Si/Cl/Ca`	0	five-stage registry elements (gas-phase; Si and especially Ca heavily depleted, Cl nearly undepleted — Orion-like 3e-6 / 3e-7 / 2e-8); adds [Si II] 34.8 μm, Si III/IV, [Cl II/III/IV] (5518/5538 density pair), Ca II, [Ca V]
`ion_metal_abs`	T	metal photoionization absorption in the band opacity (active with `use_metals`); the only gas opacity in the FUV bins

parameter	default	meaning
`emis_output`	F	write `_emis` (HDF5/FITS): leaf-by-leaf line emissivities + the state needed for maps (Section 3.)
`hei_metastable`	F	write `_heimeta` (HDF5/FITS): the He I 2^3S metastable density n_3, its photoionization rate Φ_3, plus $n_{\rm H}$, n_e, T_e, and leaf geometry (the 10830 Å absorption diagnostic; Section 3.). With `add_fuv` and `efuv_min`=4.78 the soft-UV part of Φ_3 is included
`h_coll_effects`	F	add the collisional H I Balmer component (excitation of neutral H) as separate `hc` rows/blocks; matters at fronts and hot cells
`fe_levels_full`	F	load the expanded Fe II/III models (`nlevel_fe_{2,3}_full.txt`, 16/34 levels) instead of the compact 13/14 for iron-line studies
`diffuse_field`	F	explicit ground-recombination packets
`hei_diffuse`	F	fourth diffuse channel (needs `diffuse_field`): He I excited-level recombination photons (19.82 eV, 584 Å, two-photon) — correct energy bookkeeping; moves Te by $\sim$1% at HII40, He split unchanged
`gaunt_vh14`	F	free–free Gaunt factors from the van Hoof et al. (2014) table (`data/gauntff_vh14.dat`) instead of Hummer (1988); the two agree to $\lesssim$0.5% in T_e
`metal_ne`	F	PDR: electrons from the metal cascade join the n_e closure (the only electrons beyond the I-front)
`metal_heat`	F	PDR: metal photoheating in the thermal balance
`grain_pe`	F	PDR thermal package: grain photoelectric heating + recombination cooling (Bakes & Tielens 1994) from the local FUV field (needs `add_fuv`), plus H-impact [C II] 158 μm / [O I] 63 μm cooling; scaled by `pe_scale` and the leaf dust amount
`pe_scale`	1.0	photoelectric-efficiency scale factor
`use_sec_ion`	F	secondary ionization by fast photoelectrons (Shull & van Steenberg 1985): a photoelectron with kinetic energy $E_0 = h\nu - E_{\rm th}$ above 40 eV deposits only $f_{\rm heat}(x)$ of its excess as heat and the balance drives secondary H I/He I ionizations (x = ionized fraction of the H+He nuclei); $\sim$0 in fully ionized gas, matters at fronts / PDR / hard fields
`metal_freefree`	T	include the trace metals in the free–free cooling with the net ion charge $Z_{\rm ion}$ = (stage -1), not the nuclear Z; a trace ($\sim 10^{-4}$ of the H/He free–free) correctness term, on by default

2.9 I-front re-refinement

parameter	default	meaning
`refine_front`	F	rebuild the octree on the I-front once
`refine_iter`	20	at which iteration
`refine_lbase`, `refine_lmax`	5, 7	base/front levels
`refine_eps`	1e-2	front flag: $\epsilon < x_{\rm HI} < 1-\epsilon$
`refine_dx`	0.2	or a face-neighbor $x_{\rm HI}$ jump above this

2.10 Dust in the band and dust emission

parameter	default	meaning
`ion_add_dust`	F	dust absorption in the band (dusty Strömgren)
`ion_dust_kext`	”	kext table with EUV coverage (e.g. `data/kext_albedo_WD_MW_3.1_60_D03.all_2009`); required
`ion_dust_scatter`	F	sample dust scattering (HG, albedo weights)
`dust_sed`	F	SEDust stochastic dust SED at output time
`dust_model_sed`	’astrodust’	’astrodust’/’dl07’/’zubko’
`sed_workdir`	” (auto)	SEDust `sed/` tree whose dielectric tables are read relative to it; blank auto-resolves to `<executable dir>/SEDust/sed` (the in-tree copy), so it normally needs no setting
`sed_qtable`, `sed_sizedist`	in-tree defaults	optics/size-distribution tables (relative to `sed_workdir`)
`sed_NT`, `sed_Tlo`, `sed_Thi`	200, 2.7, 5e3	SEDust temperature grid
`dust_emis_bands`	unset	up to 8 wavelengths [μm] (e.g. 8, 24, 70, 160): SEDust band emissivities of each leaf → `_dustemis` (blocks `emis_dust/wl_dust` [erg s^{-1} cm^{-3} μm^{-1}]); the Python map maker projects them into dust-band images (the IR is optically thin)
`sed_pah_live`	F	$x_{\rm HII}$-dependent PAH survival in the emission: the leaf spectrum is assembled from the model channels with the PAH channel weighted by $(x_{\rm HI} + f_{\rm ion,pah}\, x_{\rm HII})$ (needs a model with a ’PAH’ channel, i.e. astrodust); turns the 8-μm map into the classic PAH ring at the ionization front

2.11 Peel-off imaging

parameter	default	meaning
`ion_peel`	F	after convergence, one extra transport pass peels direct (stellar + diffuse) and dust-scattered contributions toward the observers → `_image` (blocks `direct_ion/scatt_ion`, plus `_fuv` with `add_fuv`) [erg s^{-1} cm^{-2} sr^{-1}]
`peel_bins`	F	bin-resolved cubes `direct_cube/scatt_cube` (n_x, n_y, n_ν) in addition to the band-integrated channel images (hardness maps, synthetic EUV/FUV photometry); the third (bin) axis is labelled by the `E_bin/dE_bin` [eV] and `wl_bin` [Å] arrays written once into the same file
`save_direc0`	F	also write the unattenuated direct image (`direc0_ion/_fuv`, and `direc0_cube` with `peel_bins`): the direct source with the path extinction $e^{-\tau}$ omitted, so `direc0/direct`$= e^{+\tau}$ is the line-of-sight optical depth toward the observer
`obsx,obsy,obsz`	unset	observer positions (or directions) in code units; or use `alpha/beta/gamma` angle lists [deg]; default: one observer on $+z$ far away
`distance`	auto	observer distance (code units)
`nxim, nyim`	129, 129	image pixels
`dxim, dyim`	auto	pixel scale [deg]; auto-sized to cover the box

Peel-off also works with the external field: an external packet's direct peel uses the Lambert weight $\max(\hat{k}_{\rm obs}\cdot\hat{n}_{\rm in}, 0)/\pi$ about its entry-surface inward normal, so only the far-side surface contributes and the image is the background field seen through the medium (the cloud in silhouette); with `save_direc0`, `direc0/direct`$= e^{+\tau}$ is again the line-of-sight optical-depth map.

2.12 Output

`out_file` (base name; extension follows `file_format` = 'hdf5' or 'fits').

3. Output files

3.1 `<base>_rates.h5`

One dataset for each extension, leaf-ordered:

extension	content
`Gamma_HI/HeI/HeII`	photoionization rates [s^{-1}]
`Heat_HI/HeI/HeII`	photoheating per particle [erg/s]
`Heat_dust`	grain heating [erg s^{-1} cm^{-3}] (when `ion_add_dust`)
`T_dust`	equilibrium mixture dust temperature [K]
`x_HI`, `x_HeI`, `x_HeII`, `n_e`, `T_e`	converged state
`x_<el>_stages`	metal ion fractions (stage, leaf)
`J_nu`	band mean intensity (bin, leaf) [erg/s/cm^2/Hz/sr]
`E_bin`, `dE_bin`	band grid [eV]
`LeafXYZ`, `LeafSize`	leaf centers and full widths (code units; for slices/cutouts)

The file header also carries the convergence record of the gas iteration: `CONVERGD` (did it converge), `NITERDON` (iterations run), `FINALDX` (final max $|\Delta x_{\rm HII}|$), and `FINALDTE` (final max $|\Delta T_e|/T_e$).

3.2 Line and emissivity outputs

`<base>_lines.txt` is a text table of integrated line luminosities on the converged state, one row for each line: element, ionization stage, wavelength [Å], L [erg s^{-1}], and $L/L(\mathrm{H}\beta)$. With `emis_output` the same line set is also written leaf-by-leaf to `<base>_emis` (HDF5, or FITS with `file_format='fits'`) as emissivity blocks in erg s^{-1} cm^{-3}, where $\sum \mathrm{emis} \cdot V = L_{\rm line}$ and the block row order matches the text table. Every line comes from one of two sources: recombination-line tables (H and He, Table 1) and n-level statistical-equilibrium solves (the metals plus the optional collisional H I component, Table 2).

Table 1: Recombination-line blocks in `<base>_emis`; wavelengths in Å unless μm noted.

block	source	lines	weighting, when present
`emis_h`	Storey & Hummer (1995) H I, case set by `line_case`	Hα 6565, Hβ 4863, Hγ 4342, Hδ 4103, Paα 1.876 μm, Paβ 1.282 μm, Brγ 2.166 μm	always; $n_e\, n_{\rm HII}$
`emis_heii`	SH95 He II (case per `line_case`)	1640, 3204, 4543, 4686, 5413, 10126	only where an He III zone exists; $n_e\, n_{\rm HeIII}$
`emis_hei`	Porter et al. (2012/2013) He I, case-B collisional-radiative (density-dependent)	3889, 4026, 4471, 5876, 6678, 7065, 7281, 10830	where He II is present; $n_e\, n_{\rm HeII}$
`emis_hc`	25-level H I collisional solve	Hα/Hβ/Hγ collisional component (separate `hc` rows)	only with `h_coll_effects`; $n_{\rm HI}$

The He I Porter table is a case-B collisional-radiative total (density-dependent; it already includes collisional excitation from the 2^3S metastable), case B only; under case A a note is printed and the H I/He II tables switch to case A while He I stays case B.

The metal line blocks `emis_<el>_<stage>` follow the *n*-level solve. One block is written for every active registry ion (element abundance > 0) that has an `nlevel_<el>_<stage>` data file. Elements C, N, O, Ne, S are active by default; Ar, Mg, Fe, Si, Cl, Ca are activated by setting their `par%abund_<El>`. Only the principal lines are listed below — the full transition list of each block is in `<base>_lines.txt`.

Table 2: Metal collisional-line blocks (`emis_<el>_<stage>`); wavelengths in Å unless μm noted.

ion (block)	principal lines
`c_2` [C II]	2325 (UV multiplet), 158 μm
`c_3` C III]	977, 1907, 1909, 178 μm
`n_2` [N II]	3064/3071, 5756, 6550, 6585, 122 μm, 205 μm
`n_3` [N III]	1750 (UV multiplet), 57 μm
`o_1` [O I]	5579, 6302, 6366, 63 μm, 145 μm
`o_2` [O II]	2471, 3727, 3730, 7321/7331
`o_3` [O III]	1661/1666, 2322/2332, 4364, 4960, 5008, 52 μm, 88 μm
`ne_2` [Ne II]	12.8 μm
`ne_3` [Ne III]	3343, 3870, 3969, 15.55 μm, 36 μm
`s_2` [S II]	4070/4078, 6718, 6733, 1.03 μm, 315 μm
`s_3` [S III]	3722, 6314, 8831, 9070, 9532, 12 μm, 18.7 μm, 33.5 μm
Optional — activate with `par%abund_<El>`	
`ar_3` [Ar III]	3006/3110, 5193, 7138, 7753, 8.99 μm, 21.8 μm
`ar_4` [Ar IV]	2854/2869, 4713, 4741, 7170 (multiplet), 56 μm, 78 μm
`mg_2` Mg II	2796, 2804
`fe_2` [Fe II]	1.26/1.32 μm, 5.34 μm, 17.9 μm, 26 μm (40 lines total)
`fe_3` [Fe III]	4659, 4703, 4881, 5011, 22.9 μm (38 lines; 16/34-level with `fe_levels_full`)
`si_2` [Si II]	2335 (multiplet), 34.8 μm
`si_3` Si III]	1206, 1883, 1892, 38 μm
`si_4` Si IV	1394, 1403
`cl_2` [Cl II]	3587/3679, 6164, 8581, 9126, 14.4 μm, 33 μm
`cl_3` [Cl III]	3344/3354, 5519, 5539, 8480 (multiplet)
`cl_4` [Cl IV]	3120/3205, 5325, 7263, 7533, 11.8 μm, 20.3 μm
`ca_2` Ca II	3935, 3970, 7293/7326, 8500/8544/8665
`ca_5` [Ca V]	3999, 5311, 6088, 3.05 μm, 4.16 μm, 11.5 μm

Some registry ions are tracked in the ionization balance but carry no line block because CHIANTI v11 has no level/collision data for them: Ar II, Cl V, Ca III, Ca IV.

State blocks. Beyond the emissivities, <base>_emis carries the state a map needs without the rates file: LeafXYZ (leaf centers) with the DIST_CM keyword, LeafSize (leaf edge length; V = (LeafSize·DIST_CM)^3), n_H, n_e, T_e, x_HI, x_HeI, x_HeII, and the x_<el>_stages metal-fraction blocks (same layout as the rates file). The Python EmisData reader (tools/python/mochii_output.py) turns these blocks into line maps without the rates file.

3.3 Spectral and image outputs

file	content
_nebcont.txt	nebular continuum L_ν (fb+ff tables, Milne, two-photon)
_dustir.txt	equilibrium-mixture dust IR spectrum (1–1000 μm) with the energy-closure header
_dustsed.txt	SEDust stochastic dust SED with PAH bands (when dust_sed)
_dustemis.h5	(with dust_emis_bands) SEDust band emissivities of each leaf, same block layout as _emis — readable by EmisData for dust-band maps
_heimeta.h5	(with hei_metastable) the He I 2^3S diagnostic: n_2s3 [cm^{-3}] (with the DIST_CM and A31 keywords), Phi_3 [s^{-1}], n_H, n_e, T_e, LeafXYZ, LeafSize — read by the Python HeiMetaData class for column and 10830 Å optical-depth maps, and convertible to a LaRT scattering-medium input by heimeta_to_lart.py
_image.h5	(with ion_peel) peel-off images for each observer: direct_ion/scatt_ion (+_fuv) [erg s^{-1} cm^{-2} sr^{-1}], suffix _obs<k> when several observers; with peel_bins, direct_cube/scatt_cube (n_x, n_y, n_ν) and the E_bin/dE_bin/wl_bin bin-grid arrays; with save_direc0, the unattenuated direc0_ion/_fuv (or direc0_cube)
_slab_Imu.txt	(with source_geometry = 'slab') the emergent intensity $I(\mu)$ at the two z-boundaries: columns μ, $I_{\rm top}$, $I_{\rm bot}$ [erg s^{-1} cm^{-2} sr^{-1}] on slab_nmu bins ($F = 2\pi\int I\mu\,\mathrm{d}\mu$), with the reflected / transmitted / absorbed energy budget (fractions of the entering L) in the header

3.4 Reading outputs in Python

tools/python/mochii_output.py is an importable module (works directly in a notebook) that reads every section file (HDF5 or FITS; read_sections), the line table (read_lines_table), and the text spectra (read_spectrum), and builds 2D maps from the emissivity file:

```
import sys; sys.path.insert(0, ".../MoCHII/tools/python")
import matplotlib.pyplot as plt
from mochii_output import EmisData

d = EmisData("run_emis.h5")
d.info()                          # blocks, fields, line count
```

```
d.line_list()                       # every stored line (block, wl)

# ready-made figures (notebook): returns (fig, ax)
fig, ax = d.plot_line("o_3", 5007.)                  # log S map
fig, ax = d.plot_line("h", 6563., unit="intensity",
                      photons=True)  # photons/s/cm^2/sr
fig, ax = d.plot_field("T_e", weight="EM")
fig, axs = plt.subplots(1, 2); d.plot_line("s_2", 6717., ax=axs[0]) \
    ; d.plot_line("s_2", 6731., ax=axs[1])   # doublet side by side

# raw arrays, when the figure is yours to build
img, ext = d.line_map("o_3", 5007., axis="z", npix=512)
img, ext = d.line_map("h", 6563., unit="intensity", photons=True)
img, ext = d.field_map("T_e", weight="EM")   # LOS-weighted average
Q = d.line_luminosity("h", 4861., photons=True)  # photon rate [1/s]
```

Line maps are flux-conserving: each leaf's luminosity is deposited with exact area overlap onto the pixel grid (AMR leaf sizes handled), so $\sum SA_{\rm pix} = L_{\rm line}$ to machine precision. `unit='flux'` (default) gives surface brightness [erg s^{-1} cm^{-2}]; `unit='intensity'` gives specific intensity [erg s^{-1} cm^{-2} sr^{-1}] $= S/4\pi$ (isotropic emission, optically thin); `photons=True` converts either to photon units (divides by hc/λ of the line). A `slab=(lo,hi)` cut restricts the line of sight. Field maps average `T_e`, `n_e`, `x_HI`, … along the line of sight with weight `'EM'` ($n_e^2 V$), `'ne'`, `'nH'`, or `'V'`. The same module doubles as a command-line tool for quick looks:

```
python3 mochii_output.py run_emis.h5                  # list lines
python3 mochii_output.py run_emis.h5 --line o_3 5007 \
        --npix 512 --log --out o3.png
python3 mochii_output.py run_emis.h5 --line h 6563 \
        --unit intensity --photons     # photons/s/cm^2/sr
python3 mochii_output.py run_emis.h5 --field T_e --weight EM
```

3.5 Slices and cutouts

A *slice* is the actual leaf value on a plane (the cell that contains each pixel), distinct from the line-of-sight *projection* above — a slice resolves the ionization front sharply where a projection blurs it. `tools/python/leaf_field.py` holds the backend (`leaf_slice`, the nearest-leaf fallback `leaf_slice_nn`, and `plot_slice`); `mochii_output.py` exposes it on the output files:

```
from mochii_output import EmisData, RatesData
r = RatesData("run_rates.h5")                # leaf state
fig, ax = r.plot_field_slice("T_e", axis="z", coord=0.0)
fig, ax = r.plot_field_slice("x_HI", axis="z", coord=0.0, log=True)
```

```
img, ext = r.slice_field("n_e", axis="x", coord=0.0,
                         bounds=(-2,2,-2,2))  # cutout window


d = EmisData("run_emis.h5")
fig, ax = d.plot_line_slice("o_3", 5007., axis="z")   # emissivity slice
```

Exact slices need the leaf full widths: the `_emis.h5` file always carries `LeafSize`, and `_rates.h5` does too; a rates file written before `LeafSize` existed falls back to a nearest-leaf slice (a one-time note is printed). The command line takes `--slice FIELD --axis z --coord 0` on either file.

3.6 He I 2^3S metastable and 10830 Å absorption

The `_heimeta` file (`hei_metastable`) is read by the `HeiMetaData` class, which adds a line-of-sight column map and the 10830 Å line-center optical depth:

```
from mochii_output import HeiMetaData
d = HeiMetaData("run_heimeta.h5")
d.info()                                  # n_2s3, Phi_3, line constants
img, ext = d.column_map("n_2s3")          # LOS column [cm^-2]
img, ext = d.tau10830_map(v_turb=0.0,     # tau_0 map (km/s turbulence)
                          component="doublet")  # blended J'=1,2 pair
fig, ax = d.plot_column("n_2s3", log=True)
fig, ax = d.plot_tau10830(component="all")
```

`column_map` is the same flux-conserving deposit as the line maps with weight nV, so $\sum NA_{\rm pix} = \sum nV$ to machine precision. `tau10830_map` deposits the line-center opacity $\kappa_0 = \frac{\sqrt{\pi}e^2}{m_e c} f \frac{\lambda}{b} n_3$, $b = \sqrt{2kT_e/m_{\rm He} + v_{\rm turb}^2}$, so a pixel value is $\int \kappa_0 \mathrm{d}l$. `component` selects the resolved $J' = 0$ line (10832.06 Å), the blended $J' = 1,2$ `'doublet'` (10833.22/10833.31 Å, 0.089 Å apart, inside the thermal Doppler width), or `'all'`; the constants (λ, f, A) come from `data/atomic/hei_metastable.txt`. The map warns when line-center $\tau_0 > 1$: the diagnostic assumes optically thin 10830 (line transfer is out of scope). The command line takes `--column n_2s3` or `--tau10830 --component doublet --vturb 5`.

Where line transfer *is* wanted (resonant-scattering emission profiles from the metastable medium), `heimeta_to_lart.py` (in `tools/python/`) converts a `_heimeta` file into a scattering-medium input for the resonance-line code LaRT (line `HeI_10833`): `--format car` (default, uniform grids) writes the all-in-one Cartesian file with `gasDen` = n_3, `T` = T_e, and zero velocities; `--format amr` writes the generic AMR table with n_3 in the `n_ion` column. `--field` selects a different leaf field for the scatterer. A worked end-to-end case (converted grid + LaRT namelist) is in `examples/hei_metastable/lart/`.

3.7 Building AMR grids

`tools/python/AMR_grid/` builds the generic AMR octree file the code reads (`grid_type='amr'`): `AMR_grid.py` (the `AMRGrid` class — uniform, density-gradient, or resolution refinement; `set_density`/`set_neutral_fraction`; `write` to

HDF5/FITS/`.fits.gz`/text with columns `x,y,z,level,nH[,metallicity,xHI,ndust]`) and `make_amr_sphere.py` (a sphere/shell CLI). `AMRGrid` carries the same `slice`/`plot_slice`/`cutout` so a grid can be inspected before a run.

4. Worked examples (the regression tests)

Each `tests/` directory is self-contained: grid builder or a symlinked grid, the input file, and a `check_*.py` gate script.

directory	what it demonstrates / gate
`g0_gamma`	rate integrals vs. analytic attenuation (bias +0.004%)
`g1_stromgren`	Strömgren sphere, case B fixed T_e; AMR $\equiv$ uniform
`g2a_thermal`	thermal H/He nebula vs. the 1D exact reference
`g2_hii`	MOCASSIN Lexington HII20/HII40 benchmarks + line fluxes
`g4_refine`	I-front re-refinement vs. native level 7
`g4_tng`	Illustris-TNG post-processing demo (shadows, maps)
`d_dusty`	dusty Strömgren; scattering bracket; T_d/IR
`peel`	peel-off images vs. the optically thin analytic flux (+0.017%); bin cubes
`uni_dda`	uniform grid: shared walk $\equiv$ octree (bit-identical); DDA walk to rounding
`pdr`	FUV/PDR switches: carbon-electron shell, photoelectric T_e
`ext_field`	isotropic external field: interior $J = J_{\rm ext}$ to 0.2% (rec and sph)
`multi_src`	multiple point sources: superposition, spectrum resolution, one-file spectra
`peel_ext`	external-field peel: silhouette flux vs. $JA_{\rm proj}/d^2$, chord τ spectrum
`qmc`	quasi-random launch: Sobol generator vs. `scipy` + Owen-net balance; the fixed-state variance gate (cell-wise $\Gamma_{\rm HI}$ gain 29$\times$ at 2^{17}, 222$\times$ at 2^{20})
`hei_meta`	He I 2^3S metastable: OH18 analytic limits, module n_3 vs. the Python closed form, Φ_3 from the band tally, and the 10830 Å column/optical-depth diagnostic (conservation + hand integral)

A minimal thermal + metals + dust input (from `d_dusty`):

```
&parameters
 par%no_photons = 4.0e7          par%iseed = 1234
 par%grid_type = 'amr'           par%amr_file = 'amr_L6.fits'
 par%distance_unit = 'pc'        par%dust_model = 'global_dgr'
 par%use_ion_band = .true.       par%nnu_ion = 32
 par%eion_min = 13.598           par%eion_max = 100.0
 par%tstar = 4.0e4               par%luminosity = 3.178e38
 par%xHI_init = 1.0e-3           par%He_abund = 0.1
 par%gas_niter = 60              par%gas_tol = 2.0e-3
 par%solve_te = .true.           par%atomic_dir = '../../data/atomic'
 par%use_metals = .true.
 par%ion_add_dust = .true.
```

```
 par%ion_dust_kext = '../../data/kext_albedo_WD_MW_3.1_60_D03.all_2009'
 par%cext_dust = 4.868e-22       par%DGR = 1.0
 par%out_file = 'run.h5'         par%file_format = 'hdf5'
/
```

5. Adding an ion

Adding an ion touches only data:

1. `tools/fitting/fit_cooling.py`: add the ion with a T_i ladder guessed from its `.elvlc` level energies; run; check the reported max error and the overlay figure.
2. `tools/fitting/fit_nlevel.py`: add the ion with the level count needed for its diagnostics; run.
3. `tools/fitting/make_element_data.py`: add the element to `ELEMENTS` (Z, stage count) and its thresholds; charge exchange entries where sources exist; run. The CHIANTI `.rrparams`/`.drparams` types 1–3 are handled (`RR`/`RR2`/`DR`/`DR2` lines).
4. Fortran: a `par%abund_<El>` field and one entry in the `species_setup` name/abundance lists (only for a NEW element; new stages of a known element need nothing).
5. Verify: extend `verify_nlevel_pyneb.py` with a diagnostic ratio where PyNeb has data; run a smoke test and check the new lines in `_lines.txt`.

Data caveats met so far: Ar II has no CHIANTI v11 level/collision data (stage tracked without lines); Fe III collision strengths differ from PyNeb's BB14 by up to $\sim$16% in 4658/4702 (known spread for iron); charge-exchange stage labels differ between the Huang (2023) and MOCASSIN transcriptions for Mg/Fe (the Huang labels are adopted; see `make_element_data.py`).

6. Practical notes

- **Bins**: 32 bins give $\sim$2% rate-integral discretization; use 64 for sub-percent work. The `add_fuv` extension carries its own bins below `eion_min`, so the ionizing resolution is unaffected.
- **FUV and the PDR**: without `add_fuv` there is no radiation beyond the I-front and the metals there recombine toward neutral; with it, FUV penetrates the neutral gas (dust extinction only) and produces the PDR-side C II and Mg II. Metal photoheating and grain photoelectric heating are not in the thermal balance, so PDR-zone temperatures are indicative only.
- **Convergence**: use `conv_crit = 'vol'` for production — the cell-max criterion stalls at the front-cell noise floor (and, with `add_fuv`, at no-heating cells where T_e pins to `te_min`). Both measures are printed each iteration. The explicit diffuse field smooths the front and converges far better.

- **Starts**: prefer an ionized initial state; a neutral start advances the front ∼one cell per iteration.
- **kext tables**: must cover the band (the astrodust table stops at 912 Å; use the D03 table for EUV work). Tables are sorted on read (the D03 file has an out-of-order seam).
- **Re-refinement**: one event per run; the old shared-memory windows are reclaimed only at exit.
- **SEDust**: `sed_workdir` must point at a SEDust `sed/` tree; the exact stochastic solve costs ∼0.05–0.1 s for each gas leaf (distributed over ranks).